



AWS Serverless Solutions

INE Cloud Series

An INE On-Demand Course

Brooks Seahorn

AWS Instructor



+ bseahorn@ine.com



+ @brooksseahorn



+ Search “Brooks Seahorn”

Topics Covered

- + Introduction to serverless
- + AWS Lambda
 - + Best Practices & IaC
- + Amazon EventBridge
- + Amazon SNS & SQS
- + Amazon API Gateway

Topics Covered

- + Labs

- + Lab 1: Lambda
- + Lab 2: EventBridge
- + Lab 3: SNS
- + Lab 4: Capstone-Complete Solution



Serverless in AWS

INE Cloud Series

An INE On-Demand Course

Brooks Seahorn

AWS Instructor



+ bseahorn@ine.com



+ @brooksseahorn



+ Search “Brooks Seahorn”

What is serverless





Lambda function

010110001010110101001010101101
010010101011011000110110110101



Lambda function



Amazon Simple Queue
Service (Amazon SQS)

Standard: 14ms to just over 100ms
FIFO: 20ms to over 600ms



Amazon EventBridge

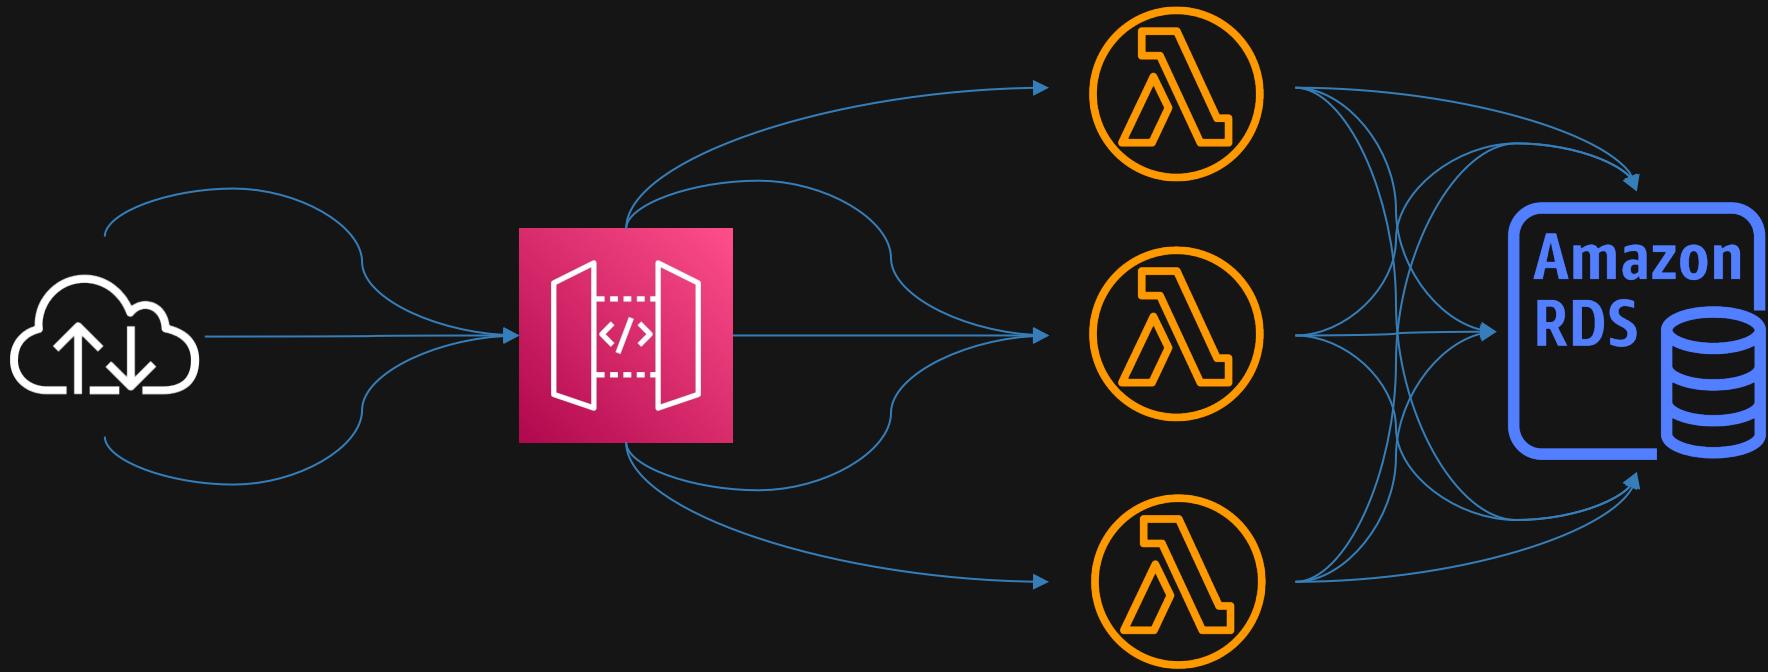
240ms to 800ms



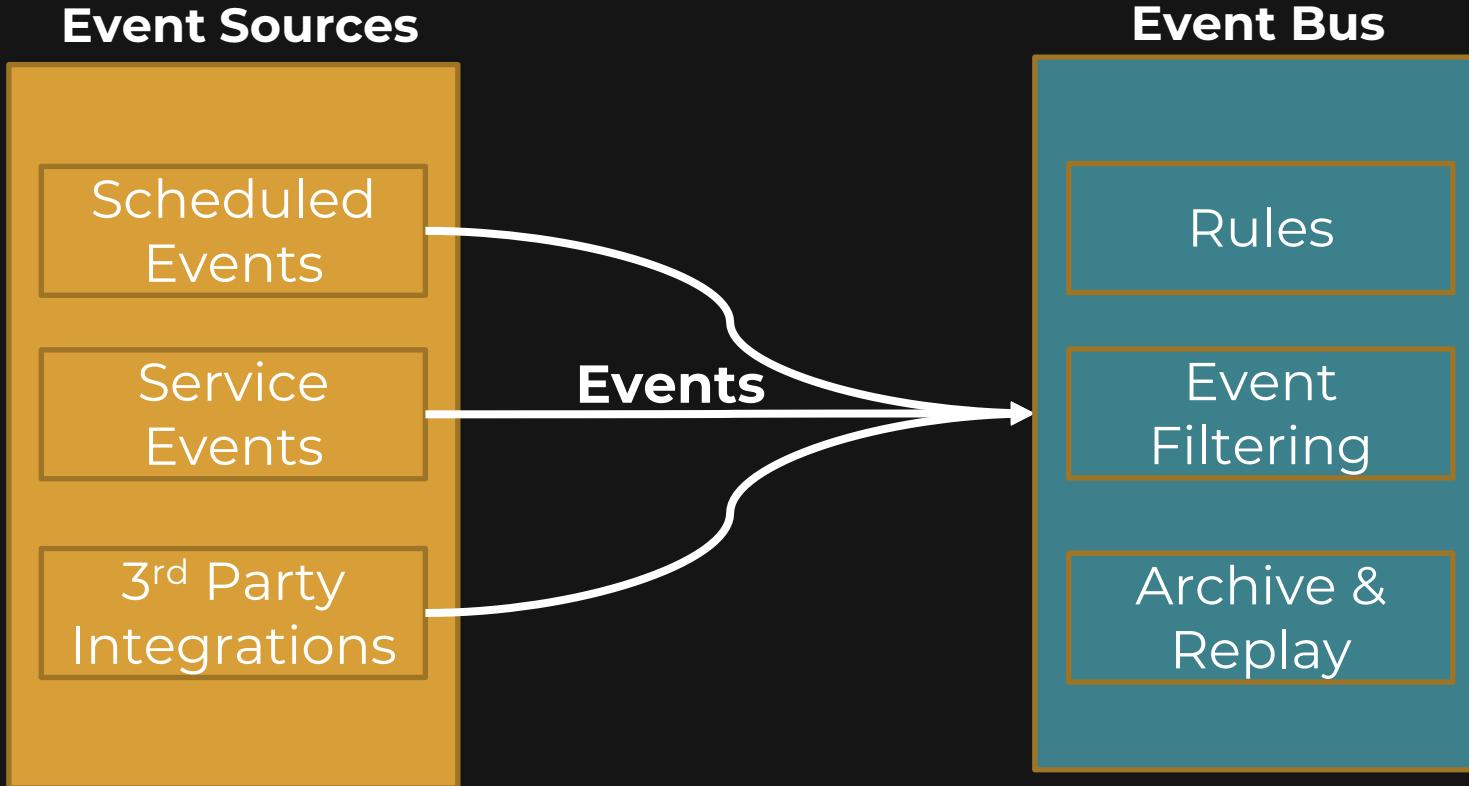
Amazon Kinesis

Standard data stream:
250ms to over 1400ms

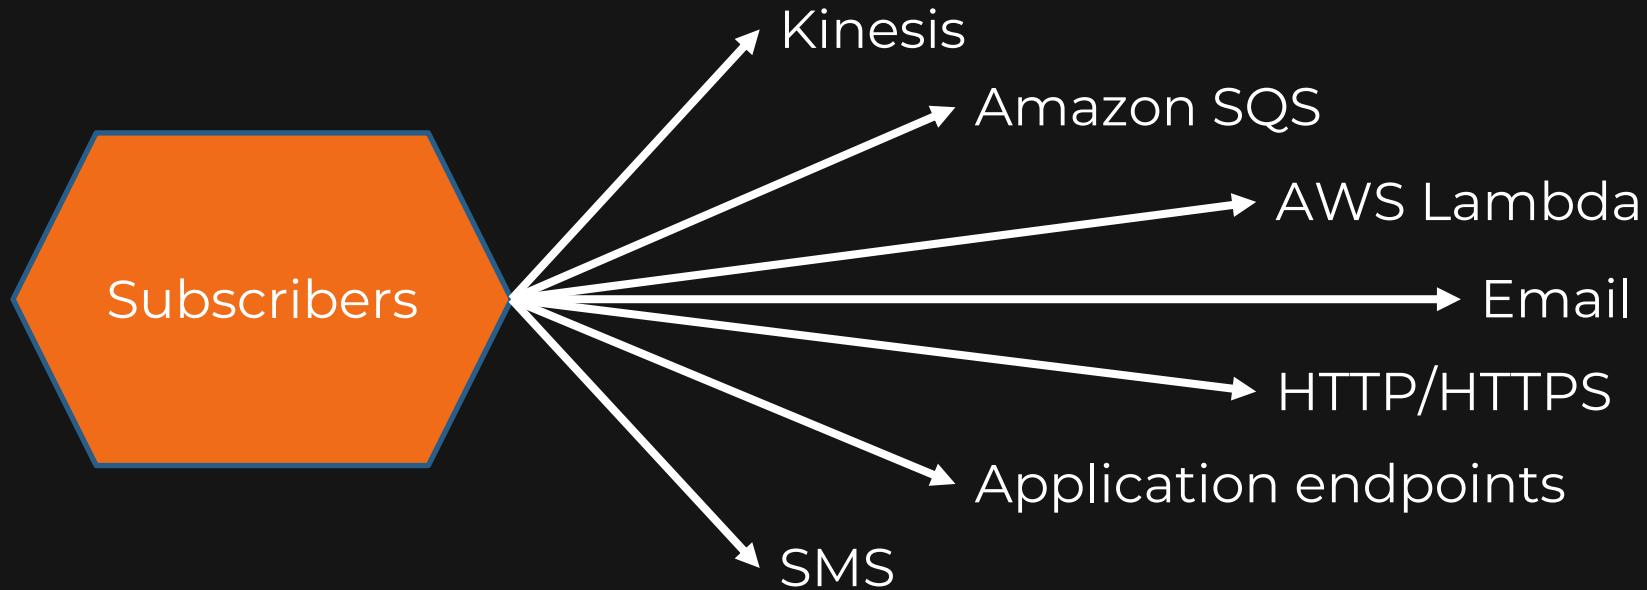
Amazon RDS Proxy



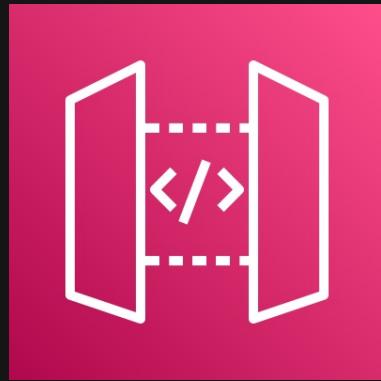
Amazon EventBridge



Subscribers

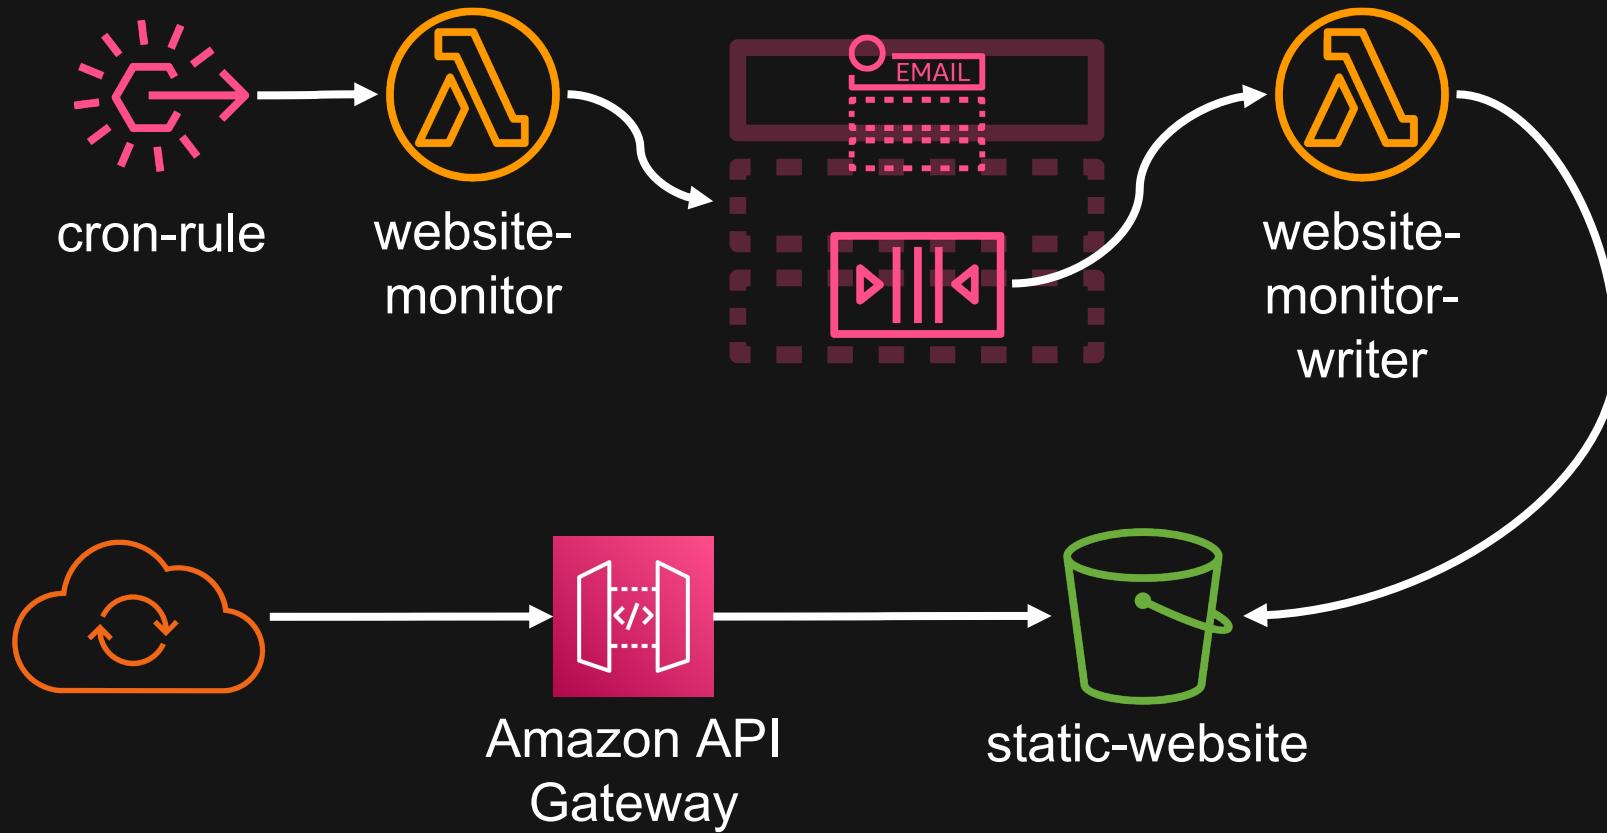


Amazon API Gateway



Amazon API
Gateway

Website Monitoring System





Serverless in AWS Conclusion

INE Cloud Series

An INE On-Demand Course

Topics Covered

- + Introduction to serverless
- + AWS Lambda
- + Amazon EventBridge
- + Amazon SNS & SQS
- + Amazon API Gateway
- + Website Monitoring Solution

Recommendations

- + Run Lab 4 again!
- + Improve on that solution
- + Review AWS' Serverless documentation
 - + <https://aws.amazon.com/serverless>
- + Review outside sources on serverless

Request from me to you

Leave a review

Thanks for watching!!!



A brief introduction to serverless

INE Cloud Series

An INE On-Demand Course

Topics

- What is serverless
- Real world serverless
- Serverless services in this course

What is serverless



What is serverless



What is serverless

```
010110001010110101001010101101  
010010101011011000110110110101
```



Real world serverless

- + Small code base
- + Limited execution time
- + Managed architecture
- + Managed statefulness

```

use num::Complex;
use std::str::FromStr;
use image::ColorType;
use image::png::PNGEncoder;
use std::fs::File;
extern crate num_cpus;
use std::env;

fn main() {
    let args: Vec<String> = env::args().collect();

    if args.len() != 5 {
        eprintln!("Usage: {} FILE PIXELS UPPERLEFT LOWERRIGHT", args[0]);
        eprintln!("Example: {} mandel.png 1000x750 -1.20,0.35 -1,0.20", args[0]);
        std::process::exit(1);
    }

    let bounds = parse_pair(&args[2], 'x')
        .expect("error parsing image dimensions");
    let upper_left = parse_complex(&args[3])
        .expect("error parsing upper left corner point");
    let lower_right = parse_complex(&args[4])
        .expect("error parsing lower right corner point");

    let mut pixels = vec![0; bounds.0 * bounds.1];

    // Commented out the following line. It is the nonconcurrent implementation.
    // render(&mut pixels, bounds, upper_left, lower_right);
    // The following is the concurrent implementation. It breaks the image
    // into horizontal bands that each are handled by a separate thread.
    let threads = num_cpus::get();
    println!("Running the set with {} thread(s).", threads);
    let rows_per_band = bounds.1 / threads + 1;
    {
        let bands: Vec<&mut [u8]> =
            pixels.chunks_mut(rows_per_band * bounds.0).collect();
        crossbeam::scope(|spawner| {
            for (i, band) in bands.into_iter().enumerate() {
                let top = rows_per_band * i;
                let height = band.len() / bounds.0;

```



```

use num::Complex;
use std::str::FromStr;
use image::ColorType;
use image::png::PNGEncoder;
use std::fs::File;
extern crate num_cpus;
use std::env;

fn main() {
    let args: Vec<String> = env::args().collect();

    if args.len() != 5 {
        eprintln!("Usage: {} FILE PIXELS UPPERLEFT LOWERRIGHT", args[0]);
        eprintln!("Example: {} mandel.png 1000x750 -1.20,0.35 -1,0.20", args[0]);
        std::process::exit(1);
    }

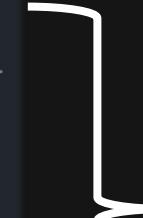
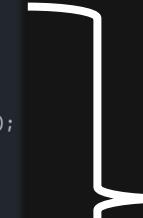
    let bounds = parse_pair(&args[2], 'x')
        .expect("error parsing image dimensions");
    let upper_left = parse_complex(&args[3])
        .expect("error parsing upper left corner point");
    let lower_right = parse_complex(&args[4])
        .expect("error parsing lower right corner point");

    let mut pixels = vec![0; bounds.0 * bounds.1];

    // Commented out the following line. It is the nonconcurrent implementation.
    // render(&mut pixels, bounds, upper_left, lower_right);
    // The following is the concurrent implementation. It breaks the image
    // into horizontal bands that each are handled by a separate thread.
    let threads = num_cpus::get();
    println!("Running the set with {} thread(s).", threads );
    let rows_per_band = bounds.1 / threads + 1;
    {

        let bands: Vec<&mut [u8]> =
            pixels.chunks_mut(rows_per_band * bounds.0).collect();
        crossbeam::scope(|spawner| {
            for (i, band) in bands.into_iter().enumerate() {
                let top = rows_per_band * i;
                let height = band.len() / bounds.0;

```



Lambda function



Lambda function



Lambda function

010110001010110101001010101101
010010101011011000110110110101



Lambda function



Lambda function



Queue



Lambda function

```

use num::Complex;
use std::str::FromStr;
use image::ColorType;
use image::png::PNGEncoder;
use std::fs::File;
extern crate num_cpus;
use std::env;

fn main() {
    let args: Vec<String> = env::args().collect();

    if args.len() != 5 {
        eprintln!("Usage: {} FILE PIXELS UPPERLEFT LOWERRIGHT", args[0]);
        eprintln!("Example: {} mandel.png 1000x750 -1.20,0.35 -1,0.20", args[0]);
        std::process::exit(1);
    }

    let bounds = parse_pair(&args[2], 'x')
        .expect("error parsing image dimensions");
    let upper_left = parse_complex(&args[3])
        .expect("error parsing upper left corner point");
    let lower_right = parse_complex(&args[4])
        .expect("error parsing lower right corner point");

    let mut pixels = vec![0; bounds.0 * bounds.1];

    // Commented out the following line. It is the nonconcurrent implementation.
    // render(&mut pixels, bounds, upper_left, lower_right);
    // The following is the concurrent implementation. It breaks the image
    // into horizontal bands that each are handled by a separate thread.
    let threads = num_cpus::get();
    println!("Running the set with {} thread(s).", threads );
    let rows_per_band = bounds.1 / threads + 1;
    {

        let bands: Vec<&mut [u8]> =
            pixels.chunks_mut(rows_per_band * bounds.0).collect();
        crossbeam::scope(|spawner| {
            for (i, band) in bands.into_iter().enumerate() {
                let top = rows_per_band * i;
                let height = band.len() / bounds.0;

```



Queue



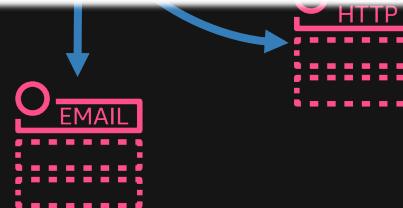
Lambda function



Increased management overhead



Increased latency





Amazon Simple Queue
Service (Amazon SQS)

Standard: 14ms to just over 100ms
FIFO: 20ms to over 600ms



Amazon EventBridge

240ms to 800ms



Amazon Kinesis

Standard data stream:
250ms to over 1400ms

Real world serverless

- + Small code base
- + Limited time execution
- + Managed architecture
- + Managed statefulness

Real world serverless

- + Serverless IS an architectural option
- + Serverless IS NOT an idealized, perfected end state

Serverless Services in AWS



AWS Lambda



Amazon EventBridge



Amazon API Gateway



Amazon Simple Notification
Service (Amazon SNS)



Amazon Simple Queue
Service (Amazon SQS)



Project: Website Monitoring System

INE Cloud Series

An INE On-Demand Course



INE

amazon.com®



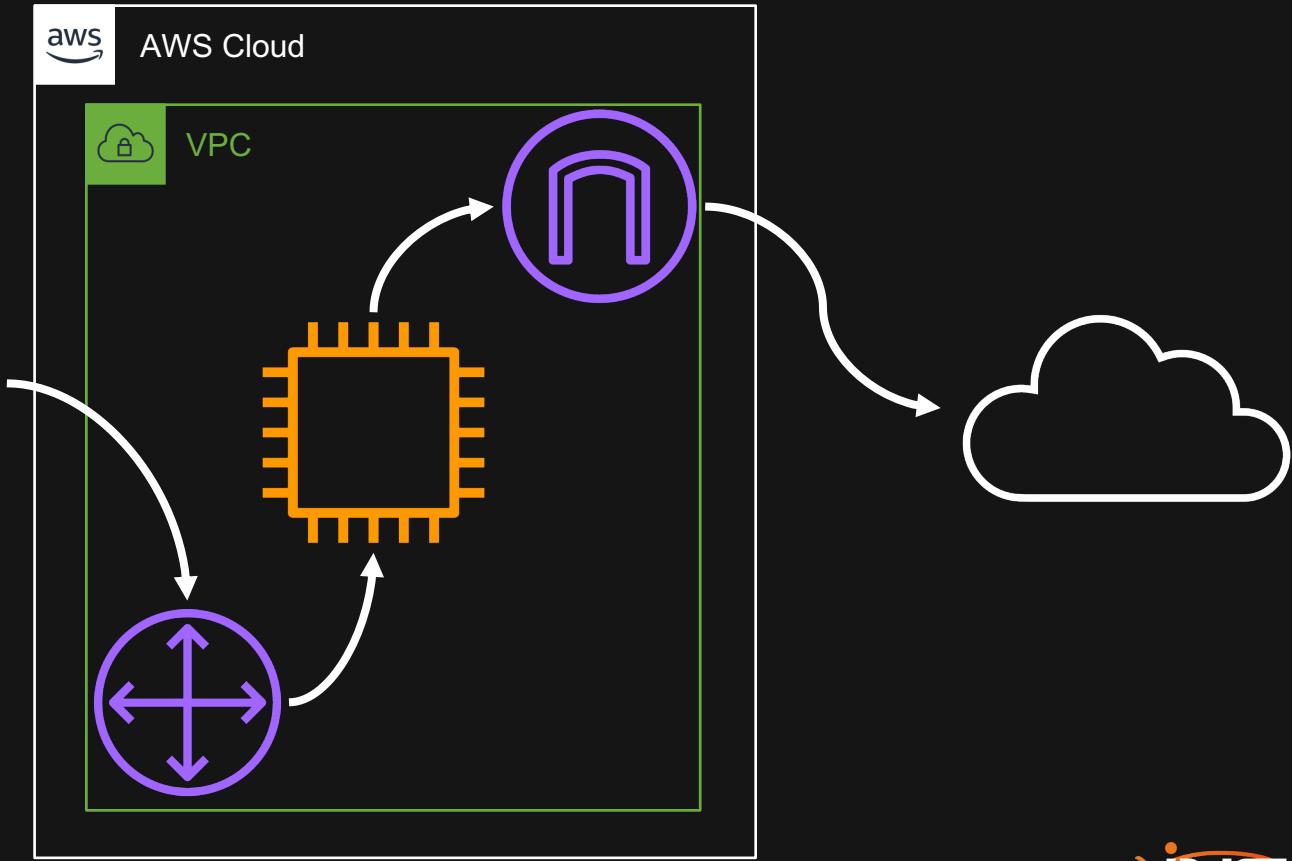
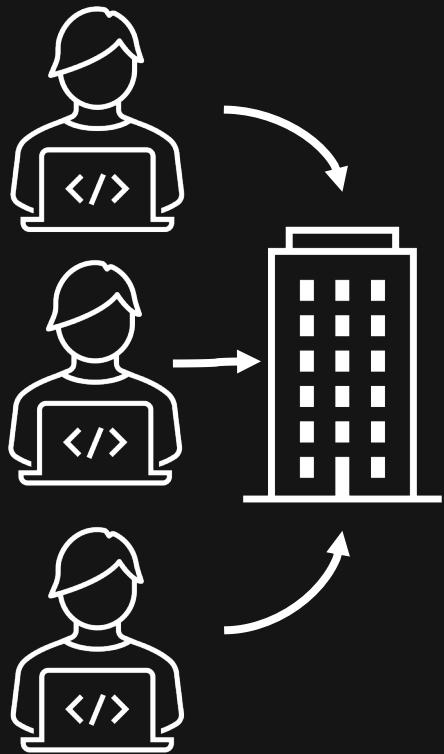
overstock™

Rakuten





Website is unreachable...

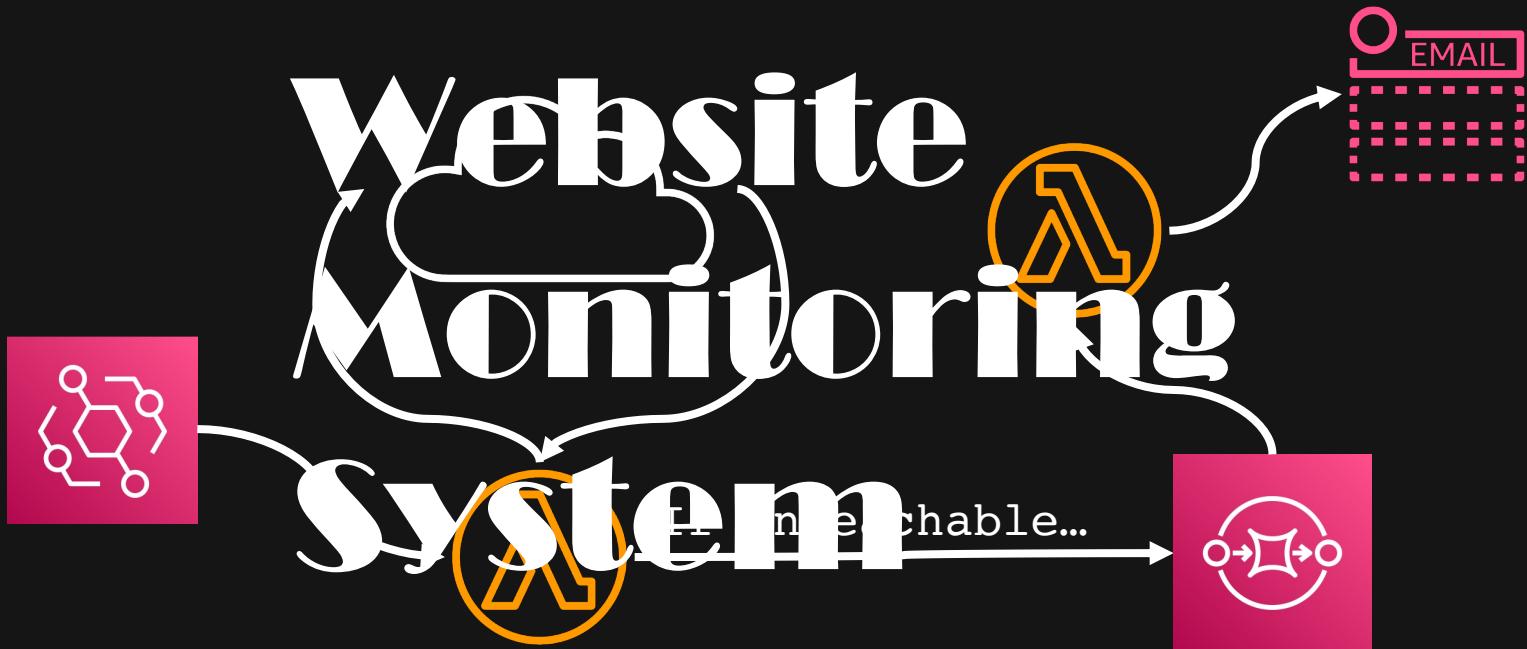


Website is unreachable...





Website Monitoring System





AWS Lambda

INE Cloud Series

An INE On-Demand Course

Topics

- AWS Lambda
- Usage
- Details
- Costs

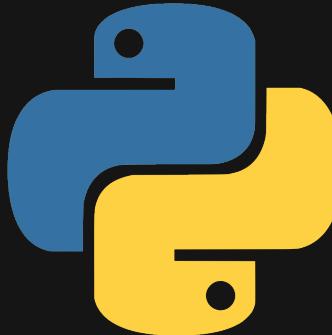
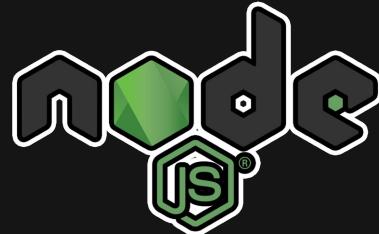
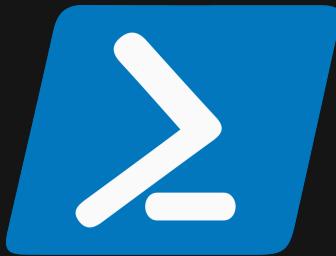
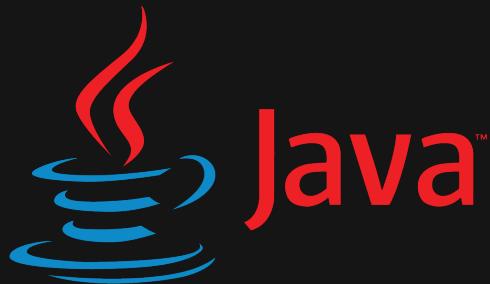
AWS Lambda Usage

- + Web applications
- + Backends (web, application, and mobile)
- + Event-driven architectures
- + Extract-transform-load

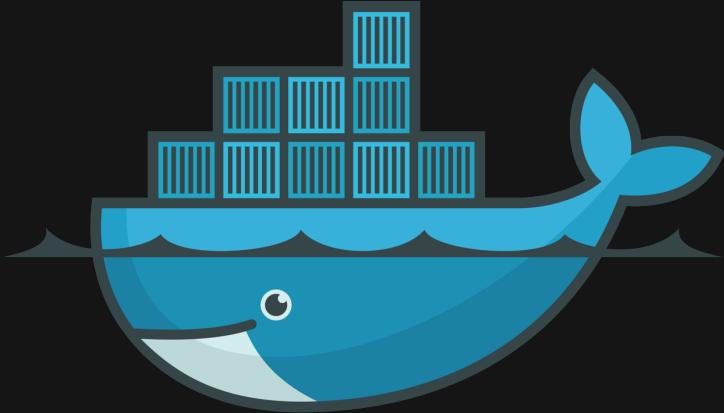
AWS Lambda

- + Serverless execution code
 - + Container based execution
 - + Managed VPC, multi-AZ deployment
- + Trigger Lambda from many AWS services
 - + Kinesis stream
 - + Write to an S3 bucket
 - + API Gateway
 - + EventBridge scheduled call
 - + Target from SQS

AWS Lambda



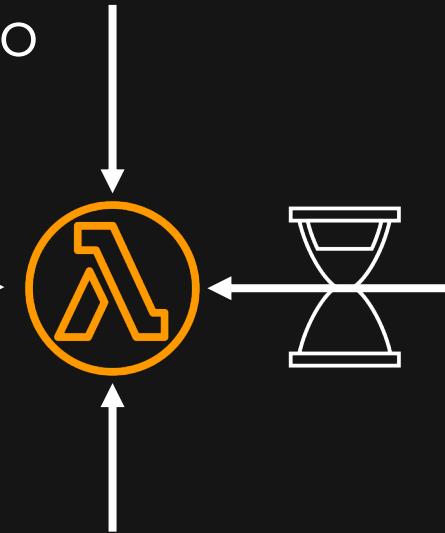
AWS Lambda



docker

AWS Lambda

- + Automatic scaling w/ no limit to capability
- + 128MB to 10,240MB Memory
- + 15-minute maximum timeout
- + Ephemeral storage
 - + 512MB to 10GB
 - + Presented as /tmp

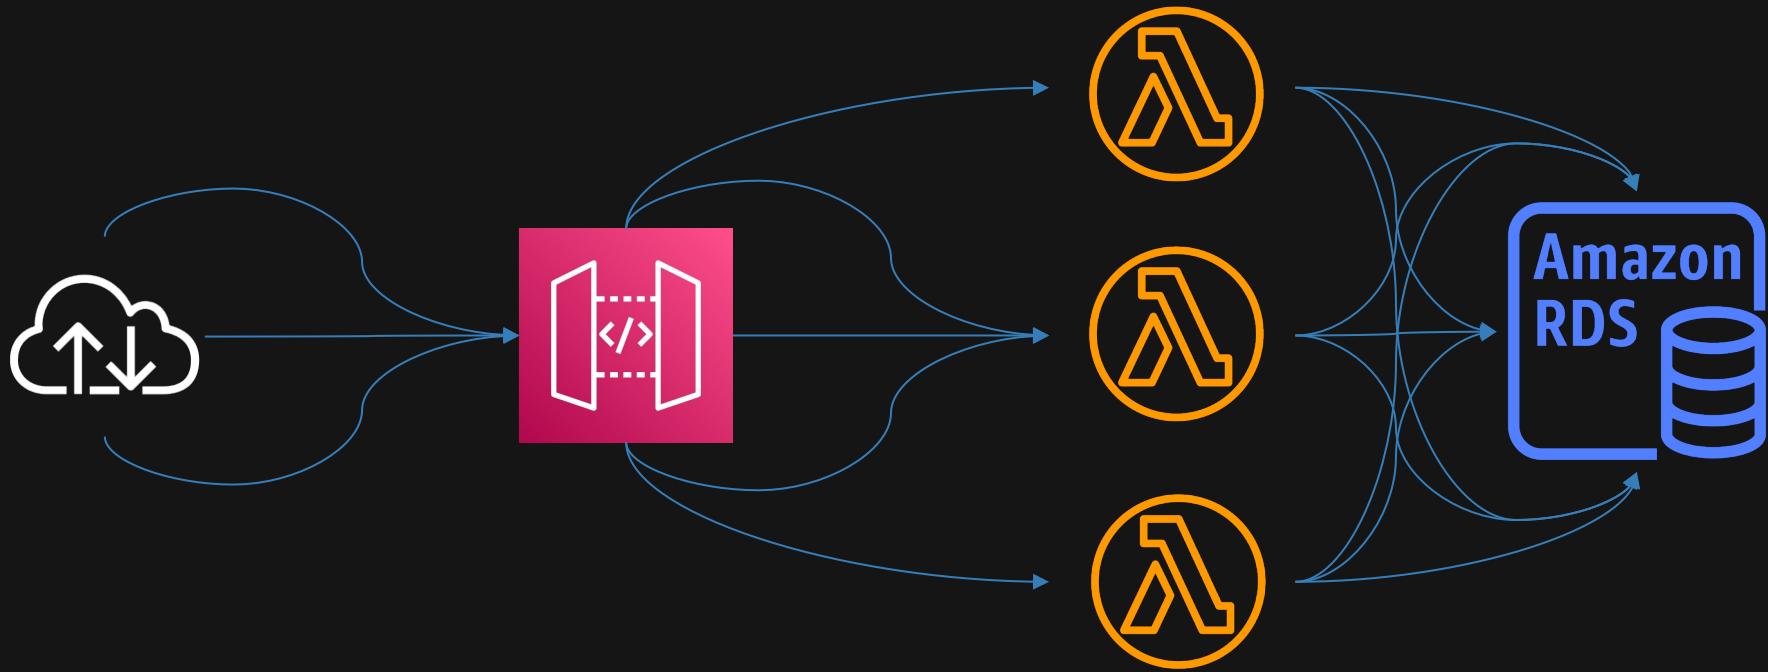


Check for ephemeral data!!!

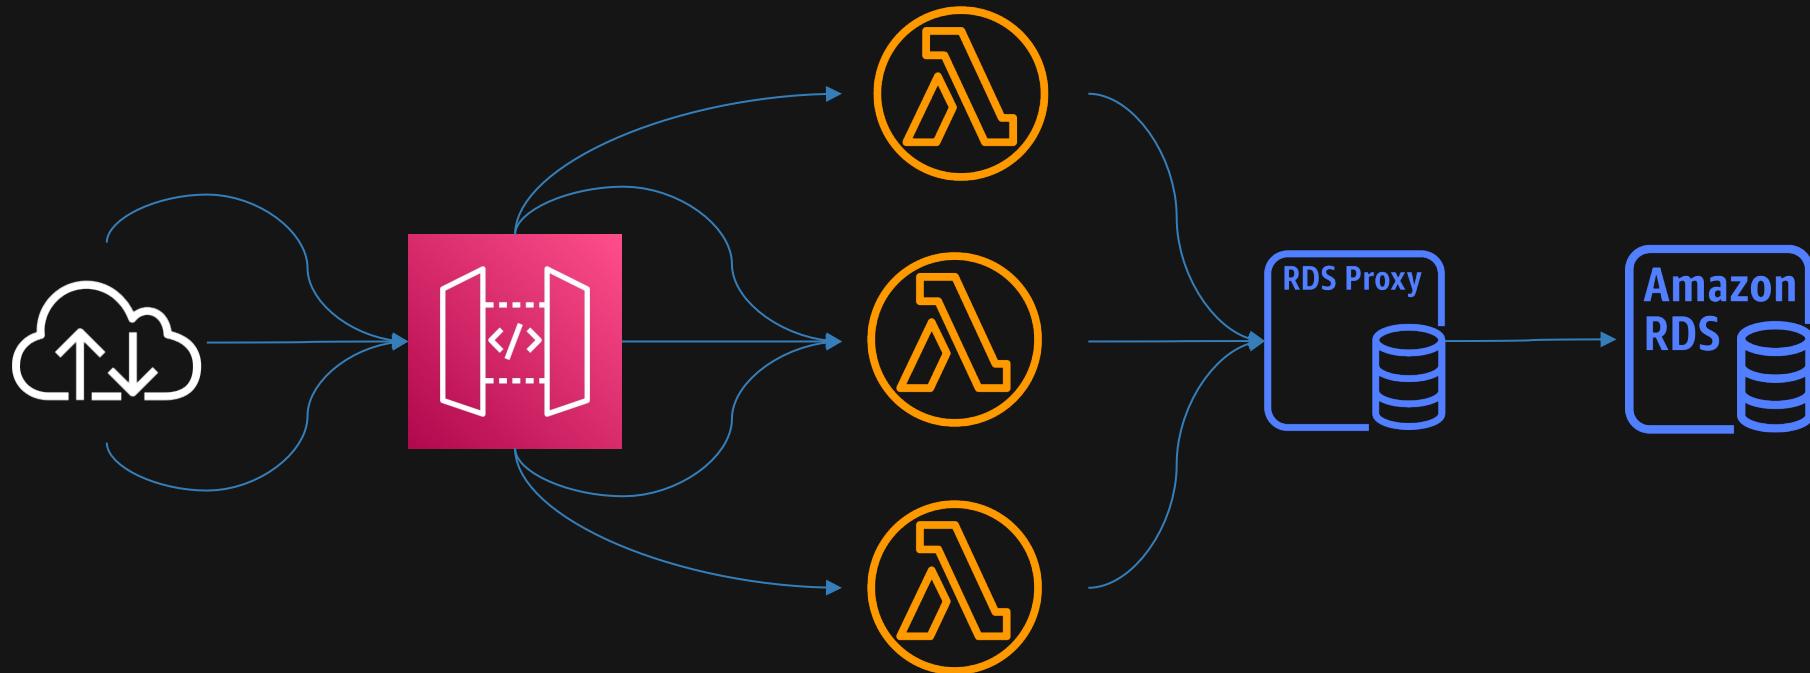
AWS Lambda

- + Automatic scaling w/ no limit to capability
- + 128MB to 10,240MB Memory
- + 15-minute maximum timeout
- + Ephemeral storage
 - + 512MB to 10GB
 - + Presented as /tmp
- + Persistent storage
 - + RDS connectivity via RDS Proxy

Amazon RDS Proxy



Amazon RDS Proxy



AWS Lambda

- + Automatic scaling w/ no limit to capability
- + Ephemeral storage
 - + 512MB to 10GB
 - + Presented as /tmp
- + Persistent storage
 - + RDS connectivity via RDS Proxy
 - + Amazon EFS

AWS Lambda Cost

- + Execution (128MB):
 - + \$0.0000000021 per 1ms
 - + First 6 Billion GB-sec/month
- + Ephemeral storage
 - + \$0.0000000309 for every GB-second

Conclusion

- + AWS Lambda
- + Usage
- + Details
- + Costs



AWS Lambda Demonstration

INE Cloud Series

An INE On-Demand Course

Demonstration

- + Web site monitor
- + Initial build
- + Add custom code
- + Create a functional test
- + Monitor with CloudWatch



AWS Lambda Best/Worst Practices

INE Cloud Series

An INE On-Demand Course

Topics

- Best practices
- Bad practices

Best Practices – Code Organization

```
21 def lambda_handler(event, context):|      You, 2 days ago • Add files via upload ...
22     SITES = [item for item in temp_sites.split(",") if item]
23
24     for site in SITES:
25         NOW = str(datetime.now())
26         print('Checking ', site)
27         try:
28             req = Request(site, headers={'User-Agent': 'AWS Lambda'})
29             if not validate(str(urlopen(req).read())):
30                 raise Exception('Validation failed')
31         except:
32             print('Check failed!')
33             raise
34         else:
35             print('Check passed!' + NOW)
36             # return NOW
37         finally:
38             print(NOW)
```

Best Practices – Code Organization

```
20  def check_site(event, context):
21      SITES = [item for item in temp_sites.split(",") if item]
22
23      for site in SITES:
24          NOW = str(datetime.now())
25          print('Checking ', site)
26          try:
27              req = Request(site, headers={'User-Agent': 'AWS Lambda'})
28              if not validate(str(urlopen(req).read())):
29                  raise Exception('Validation failed')
30          except:
31              print('Check failed!')
32              raise
33          else:
34              print('Check passed!' + NOW)
35              # return NOW
36          finally:
37              print(NOW)
38
39  def lambda_handler(event, context):
40      check_site(event, context)
```

Best Practices – Performance

Function Logs

```
START RequestId: 8d0c1ad8-1caa-4314-b5ff-d1e4d6190c40 Version: $LATEST
Checking https://www.amazon.com
Check passed!2022-11-07 18:30:08.104088
2022-11-07 18:30:08.104088
Checking https://www.ebay.com
Check passed!2022-11-07 18:30:08.291278
2022-11-07 18:30:08.291278
END RequestId: 8d0c1ad8-1caa-4314-b5ff-d1e4d6190c40
REPORT RequestId: 8d0c1ad8-1caa-4314-b5ff-d1e4d6190c40
    Duration: 875.16 ms Billed Duration: 876 ms      Memory Size: 128 MB Max Memory Used: 45 MB
```

Best Practices – Performance

Duration: 875.16 ms

Billed Duration: 876 ms

Memory Size: 128 MB

Max Memory Used: 45 MB

Best Practices – Performance

AWS Lambda Power Tuner

<https://github.com/alexcasalboni/aws-lambda-power-tuning>



Best Practices – Other

- + Reuse: Subsequent calls to your function can reuse resources
- + Use environment variables
- + Minimize deployment size
 - + Use Lambda Layers
- + Avoid duplicated event processing (idempotent code)
- + SQS visibility timeout < invocation time

Best Practices – Other

- + Use most restrictive permission
- + Use versioning with Lambdas
- + Remove unused Lambdas

Best Practices – Performance

- + CloudWatch Metrics
 - + **Invocations:** # times the function is called
 - + **Errors:** Invocation # that error
 - + **DestinationDeliveryFailures:** Lambda fails to send an event to a destination
 - + **Throttles:** Rejected invocations
 - + **Duration:** Processing time
 - + **ConcurrentExecutions:** # of function instances processing

Worst Practices

- + Creating a Lambda monolith
- + Recursion (infinite loops with Lambda)
- + Lambda to Lambda calls



Amazon Lambda – Infrastructure as Code

INE Cloud Series

An INE On-Demand Course

Topics

- Security definitions
- Lambda definition
- Archive utility



Amazon EventBridge

INE Cloud Series

An INE On-Demand Course

Topics

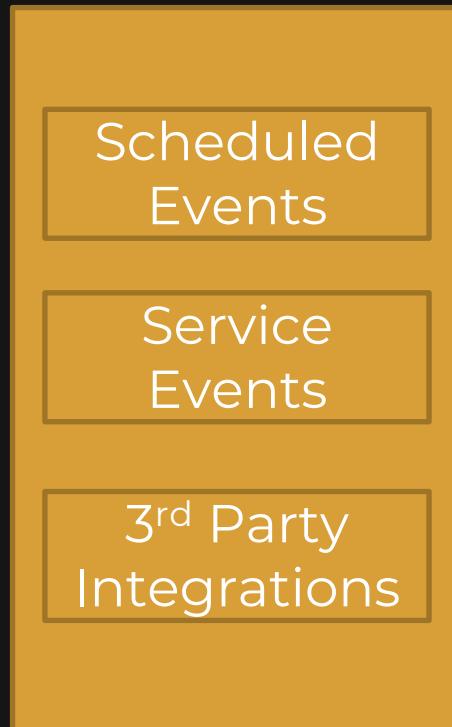
- Amazon EventBridge
- Architecture
- Demonstration

Amazon EventBridge

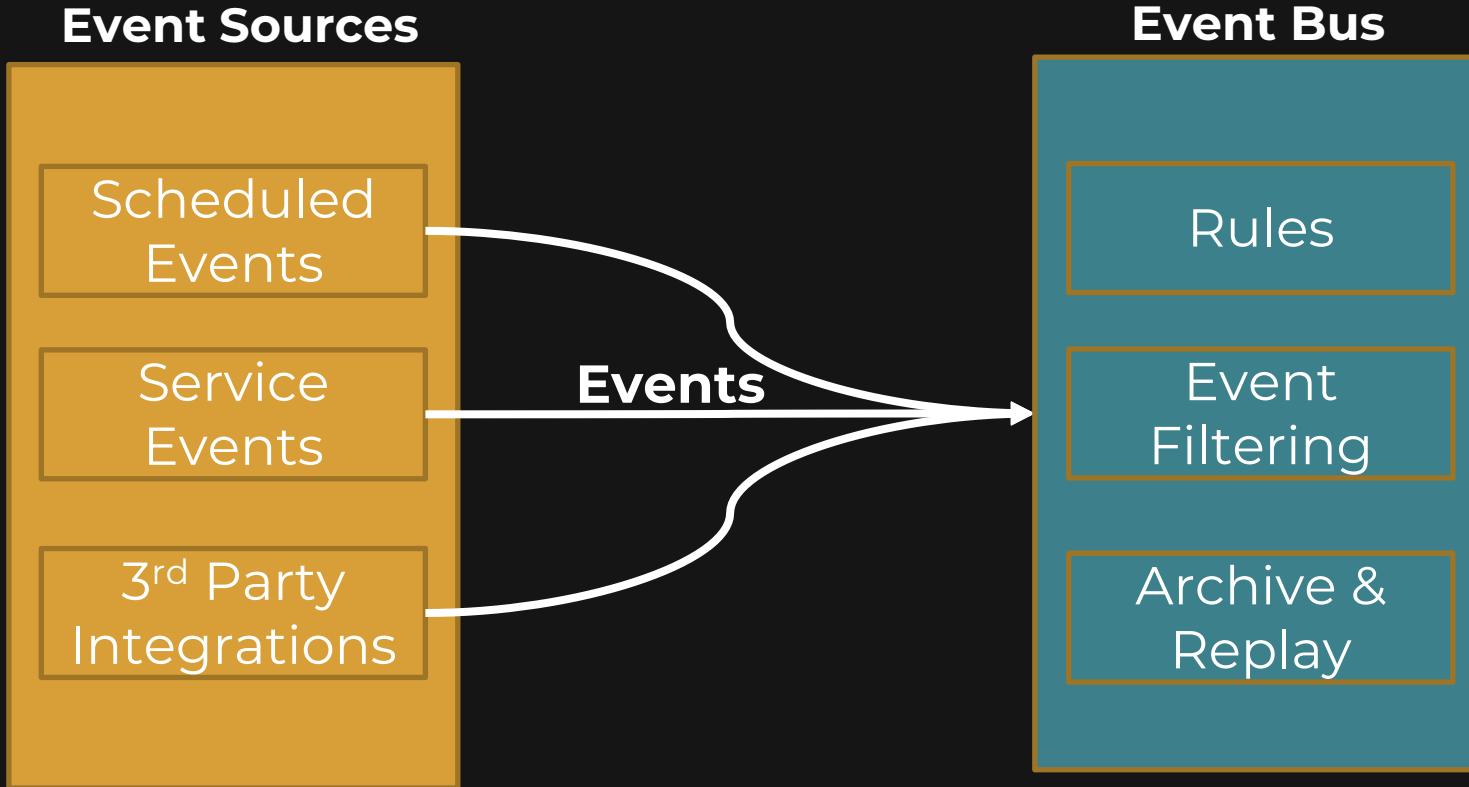
- + Build event-driven architectures
- + Connect SaaS (3rd party providers)
- + Low code, low overhead

Amazon EventBridge

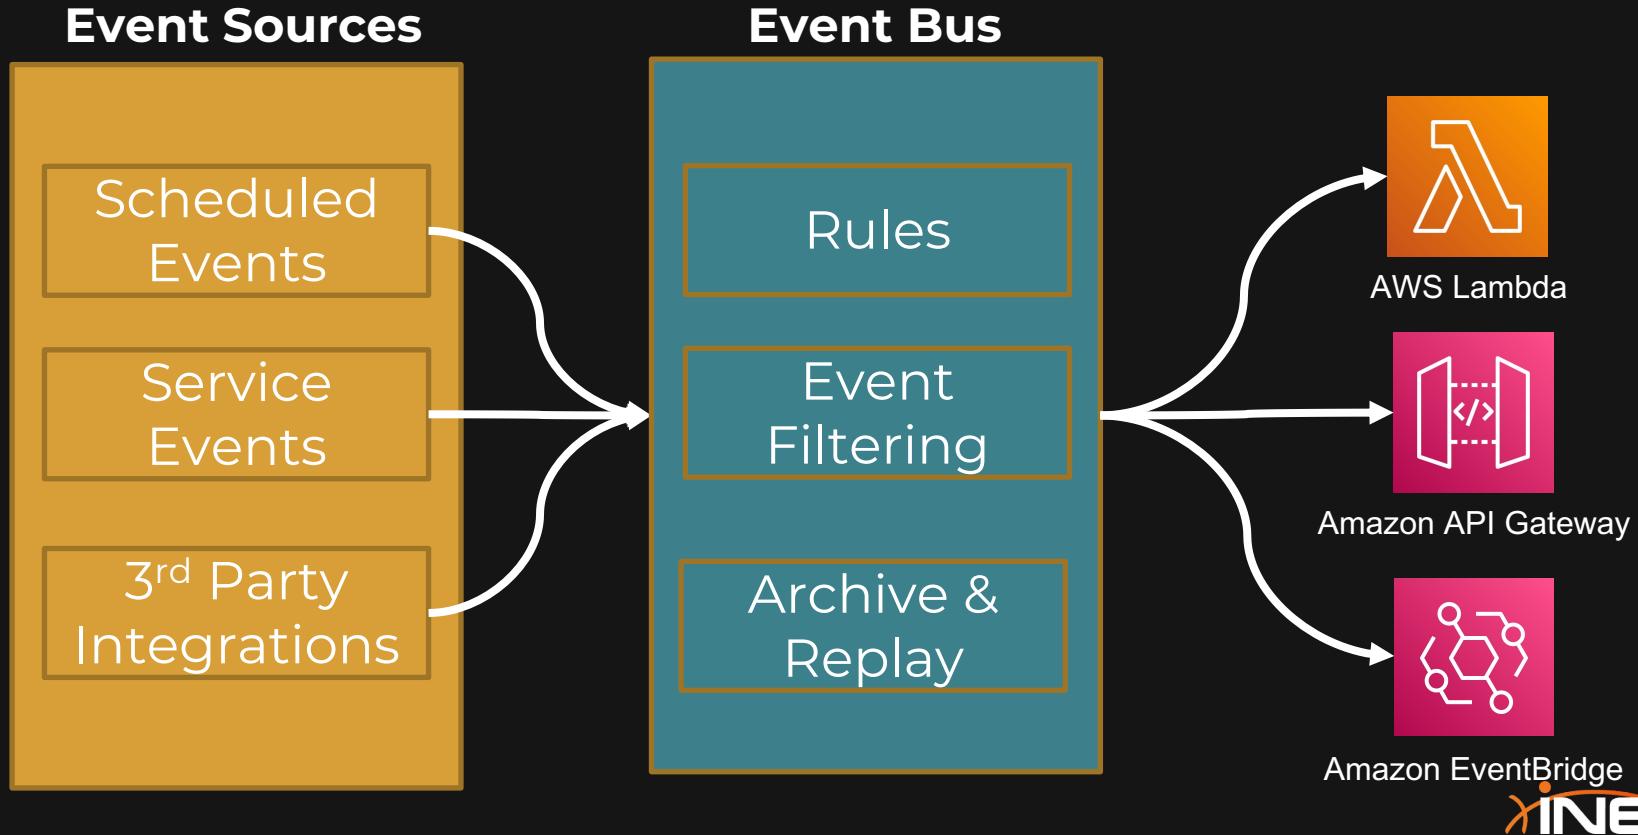
Event Sources



Amazon EventBridge



Amazon EventBridge



Amazon EventBridge



Demonstration

- + Create a rule in EventBridge
- + Call the Website Monitor on a schedule



Amazon EventBridge - Update

INE Cloud Series

An INE On-Demand Course

Topics

- November '22 Update

Amazon EventBridge

Amazon EventBridge

X

- Developer resources
 - Learn 
 - Sandbox
 - Quick starts
- Events
 - Event buses
 - Rules
 - Global endpoints
 - Archives
 - Replays
- Integration
 - Partner event sources
 - API destinations

Application Integration

Amazon EventBridge

Build event-driven applications at scale

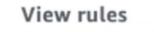
Amazon EventBridge is a serverless event bus that makes it easier to build event-driven applications at scale using events generated from your applications, integrated Software-as-a-Service (SaaS) applications, and AWS services.

How it works

 Serverless 101: Amazon EventBridge - old 

Create a new rule

Create a rule. Choose an AWS service, SaaS app or custom app as event source, define event pattern, and attach an AWS service or SaaS apps via API Destination as target(s).

Schema registry

Find existing AWS service event schemas, generate one from an event bus, or create your own custom schemas.

Automatically generate code bindings and IDE integrations for events in EventBridge.

Amazon EventBridge

Amazon EventBridge

X

Developer resources

Learn

Sandbox

Quick starts

Events

Event buses

Rules

Global endpoints

Archives

Replays

Scheduler

Schedules

Schedule group

Application Integration

Amazon EventBridge

Build event-driven applications at scale

Amazon EventBridge is a serverless event bus that makes it easier to build event-driven applications at scale using events generated from your applications, integrated Software-as-a-Service (SaaS) applications, and AWS services.

Get started

EventBridge Rule

A rule matches incoming events and sends them to targets for processing.

EventBridge Schedule

A schedule invokes a target one-time or at regular intervals defined by a cron or rate expression.

EventBridge Schema registry

Schema registries collect and organize schemas.

[Create rule](#)

Pricing



Conclusion

- + Schedules are now 1st class citizens in the EventBridge Scheduler service
- + Create services separate from the rules themselves



Amazon Simple Notification Service & Demonstration

INE Cloud Series

An INE On-Demand Course

Topics

- Amazon SNS
- Topic
- Subscriber
- Demonstration

Amazon SNS

- + Amazon Simple Notification Service
- + Pub/Sub, SMS, & mobile push
- + App-2-App (A2A) & App-2-Person (A2P)

This is a message
about something
real important.

ECS!!!

Is Python performant?

Cloud networking...where
does this stuff go?!?!?!

No no, you're crazy.
Write everything in VBA.

Mechanical
keyboards...great for Zoom
calls. ;)

Mitchell!!!



Amazon SNS

- + Topic – Message Channel
- + Subscriber – Delivery Endpoint



Topic



Subscribers

Topics

- + Message channel
- + Creating a queue: ordering, encryption, access, logging



Topic

Topics

Type [Info](#)

Topic type cannot be modified after topic is created

FIFO (first-in, first-out)

- Strictly-preserved message ordering
- Exactly-once message delivery
- High throughput, up to 300 publishes/second
- Subscription protocols: SQS

Standard

- Best-effort message ordering
- At-least once message delivery
- Highest throughput in publishes/second
- Subscription protocols: SQS, Lambda, HTTP, SMS, email, mobile application endpoints

Topics

```
→ PublishRequest pubRequest =  
    PublishRequest.builder()  
        .topicArn(topicArn)  
        .subject(subject)  
        .message(payload)  
        .messageGroupId(groupId)  
        .messageDuplicationId(dedupId)  
        .messageAttributes(attributes)  
        .build();
```

Topics

```
PublishRequest pubRequest =  
    PublishRequest.builder()  
        .topicArn(topicArn)  
        .subject(subject)  
        .message(payload)  
        .messageGroupId(groupId)  
        .messageDuplicationId(dedupId)  
        .messageAttributes(attributes)  
        .build();
```

Topics

```
PublishRequest pubRequest =  
PublishRequest.builder()  
    .topicArn(topicArn)  
    .subject(subject)  
    .message(payload)  
    .messageGroupId(groupId)  
    .messageDuplicationId(dedupId)  
    .messageAttributes(attributes)  
    .build();
```



Topics

Type [Info](#)

Topic type cannot be modified after topic is created

FIFO (first-in, first-out)

- Strictly-preserved message ordering
- Exactly-once message delivery
- High throughput, up to 300 publishes/second
- Subscription protocols: SQS

Standard

- Best-effort message ordering
- At-least once message delivery
- Highest throughput in publishes/second
- Subscription protocols: SQS, Lambda, HTTP, SMS, email, mobile application endpoints

Optional Features

- + Encryption
 - + Automatic, only on queue
- + Access policies
 - + Who can publish/subscribe
- + Data protection policies
- + Filters
 - + 200 per topic, 10,000 per account
- + HTTP/S retry policies
- + Delivery logging
- + Tags

Optional Features

FIFO

- + Encryption
- + Access policies
- + Delivery logging
- + Tags

Standard

- + Encryption
- + Access policies
- + Data protection policies
- + HTTP/S retry policies
- + Delivery logging
- + Tags

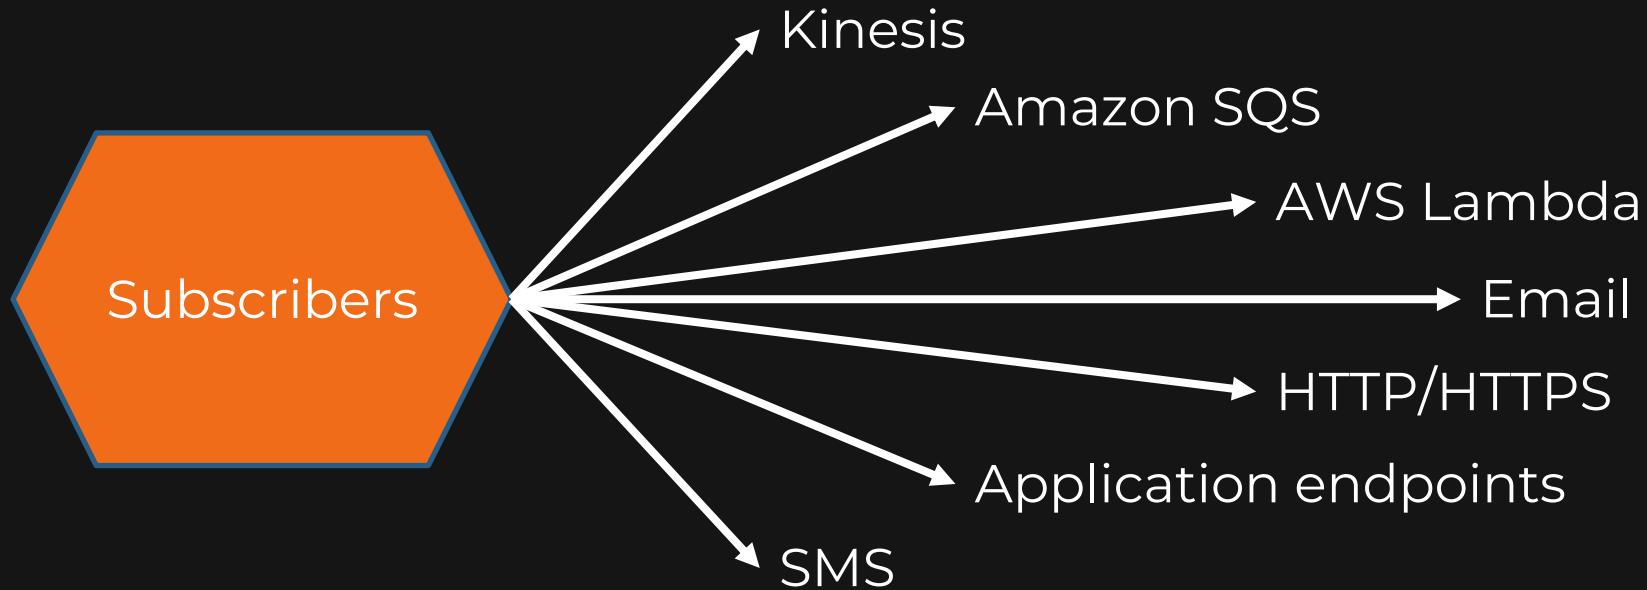
Subscribers

- + Tie topic to a delivery endpoint
- + **Fan out** to subscriber endpoints



Subscribers

Subscribers



Conclusion

- + Amazon SNS
- + A2P or A2A, Pub/Sub model
- + Topic & Subscriber
- + FIFO & Standard
- + Options



Amazon Simple Queue Service & Demonstration

INE Cloud Series

An INE On-Demand Course

Topics

- Amazon SQS
- Features
- Visibility Timeout
- Demonstration

Amazon Simple Queue Service

- + Decouple infrastructure
- + “Linkage” between processing elements

Amazon Simple Queue Service

- + Managed queue
- + Message size up to 256KB
 - + Each 64KB is billed as 1 request (*...a 256KB message is billed as 4 requests.*)
- + Up to 14-day message retention
- + Server-side encryption & dead-letter queue
 - + SSE & DLQ
- + Message “locking” to stop duplicate processing

SQS “Message locking”



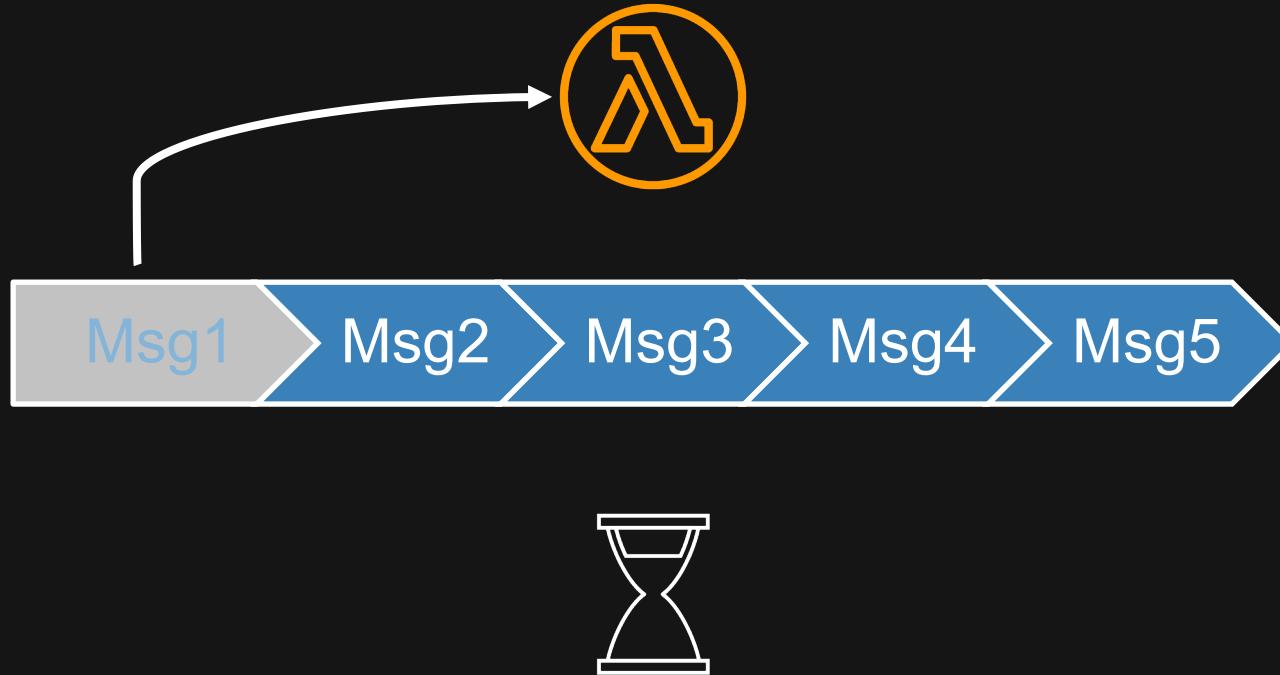
SQS “Message locking”



SQS “Message locking”

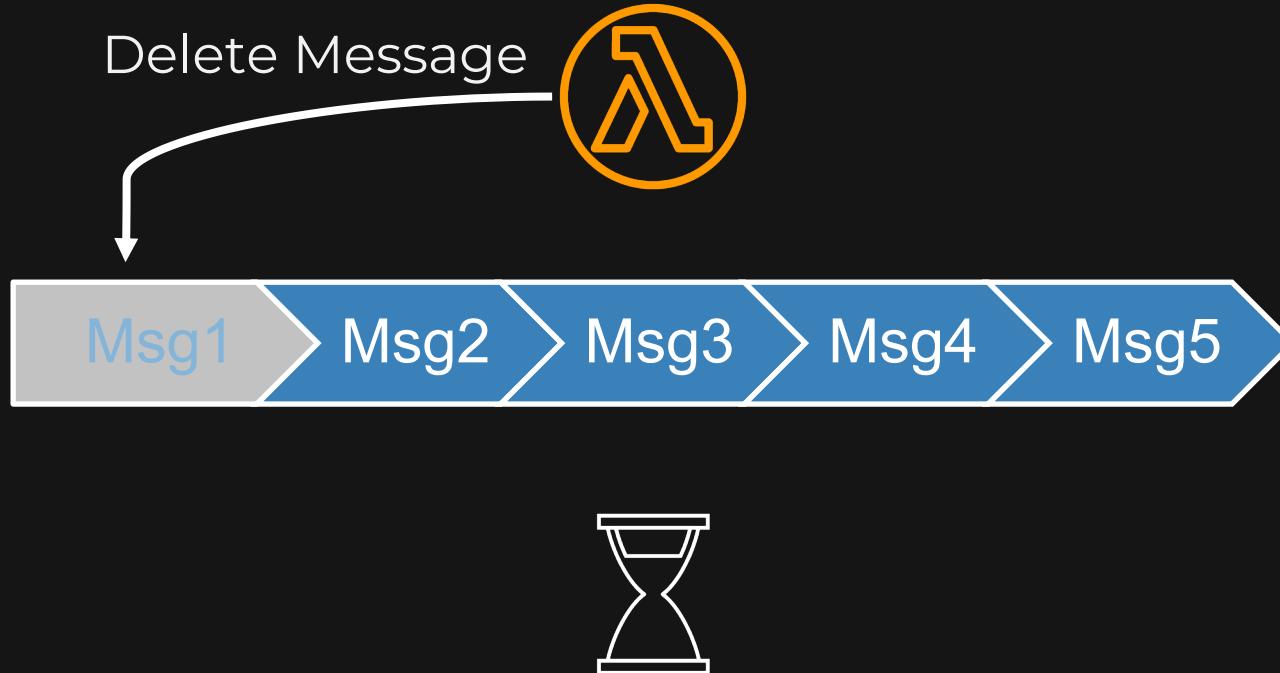


SQS “Message locking”



Visibility Timeout

SQS “Message locking”



Visibility Timeout

SQS “Message locking”



Visibility Timer

SQS Visibility Timeout



Visibility Timer

Maximum time for
processing before we retry
with another process?

Conclusion

- + Amazon SQS
- + Managed, decoupling
- + FIFO & Standard
- + Dead letter queue
- + Visibility timeout



Amazon API Gateway

INE Cloud Series

An INE On-Demand Course

Topics

- Why you need an API gateway
- Understanding API
- Amazon API Gateway

The missing link



Enable function URL [Info](#)

Use function URLs to assign HTTP(S) endpoints to your Lambda function.

The missing link



The missing link



What is an API then?



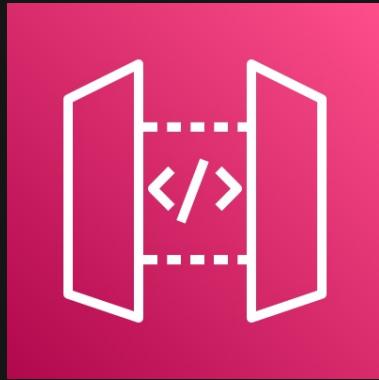
Photo by Tim Douglas : <https://www.pexels.com/photo/couple-of-barista-coworkers-standing-in-street-near-cafe-6205486/>

Photo by Andrea Piacquadio: <https://www.pexels.com/photo/shallow-focus-photo-of-woman-smiling-while-holding-pen-and-paper-3801649>

Photo by Ron Lach : <https://www.pexels.com/photo/group-of-people-sitting-by-the-table-beside-man-in-black-suit-8921559/>

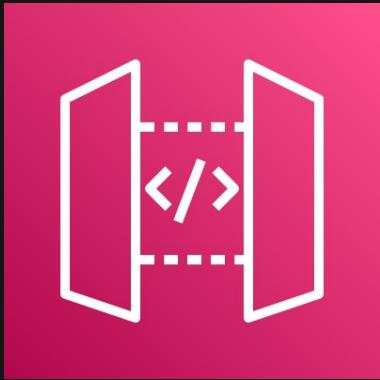
Photo by Timur Saglambilek: <https://www.pexels.com/photo/man-wearing-black-apron-near-two-silver-metal-cooking-pot-66639/>

Amazon API Gateway



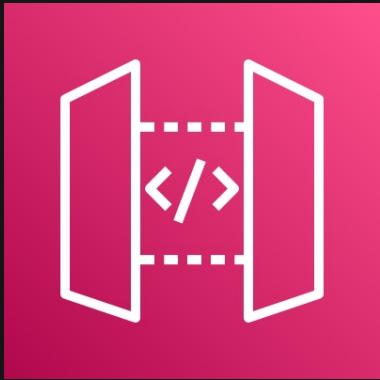
Amazon API
Gateway

Amazon API Gateway



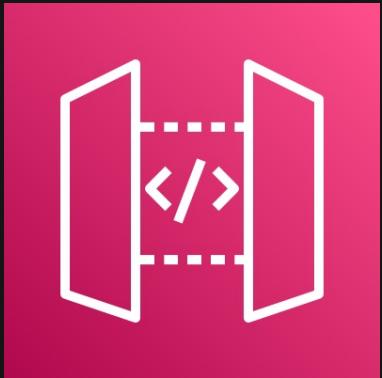
- + Managed service
- + Endpoint types
 - + HTTP API
 - + Restful API
 - + WebSocket API
 - + Real-time, 2-way communications
- + VPC Private Endpoints

Amazon API Gateway



- + Location:
 - + Regional
 - + Edge-Optimized
 - + Think **CloudFront**
 - + Private API
 - + Think **VPC**

Amazon API Gateway



- + Deploy multiple versions of your API
- + Performant w/ throttling
- + Integrates w/ CloudWatch, IAM, & Amazon Cognito
- + Cost: \$0.90/million requests

Demonstration

- + The PetStore API
 - + Understanding standard API
 - + Testing
 - + Log analysis
- + Quick build

Conclusion

- + API Gateway
- + Service particulars
- + Endpoints
- + Configuration



Serverless Lab 4: Complete Architecture Build

INE Cloud Series

An INE On-Demand Course

LAB 4 – Complete Architecture



Amazon
EventBridge

LAB 4 – Complete Architecture



cron-rule

LAB 4 – Complete Architecture



cron-rule

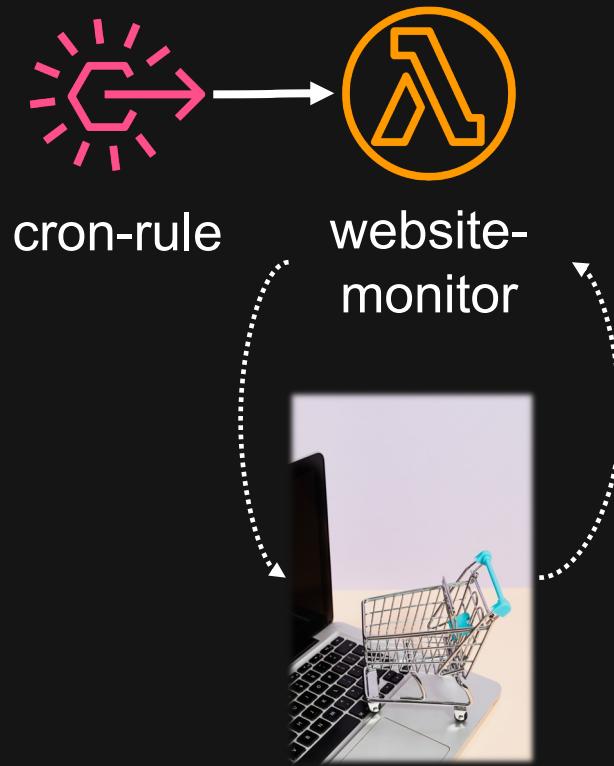


AWS Lambda

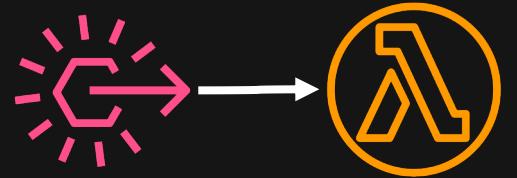
LAB 4 – Complete Architecture



LAB 4 – Complete Architecture



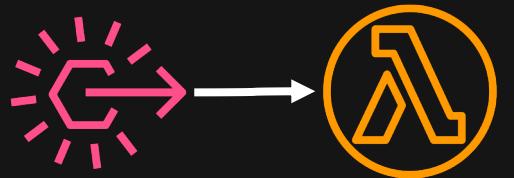
LAB 4 – Complete Architecture



cron-rule

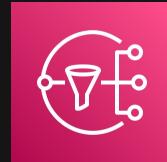
website-
monitor

LAB 4 – Complete Architecture



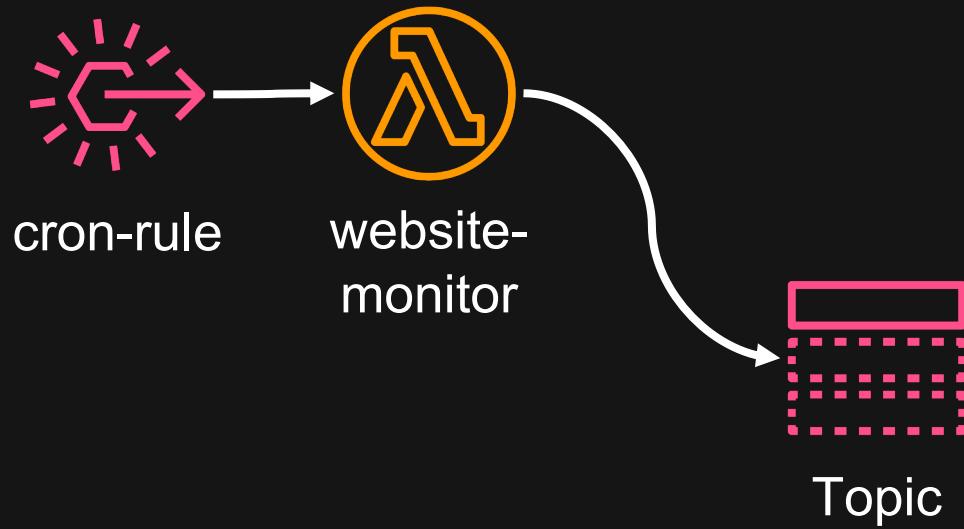
cron-rule

website-
monitor



Amazon Simple
Notification Service
(Amazon SNS)

LAB 4 – Complete Architecture



LAB 4 – Complete Architecture



LAB 4 – Complete Architecture

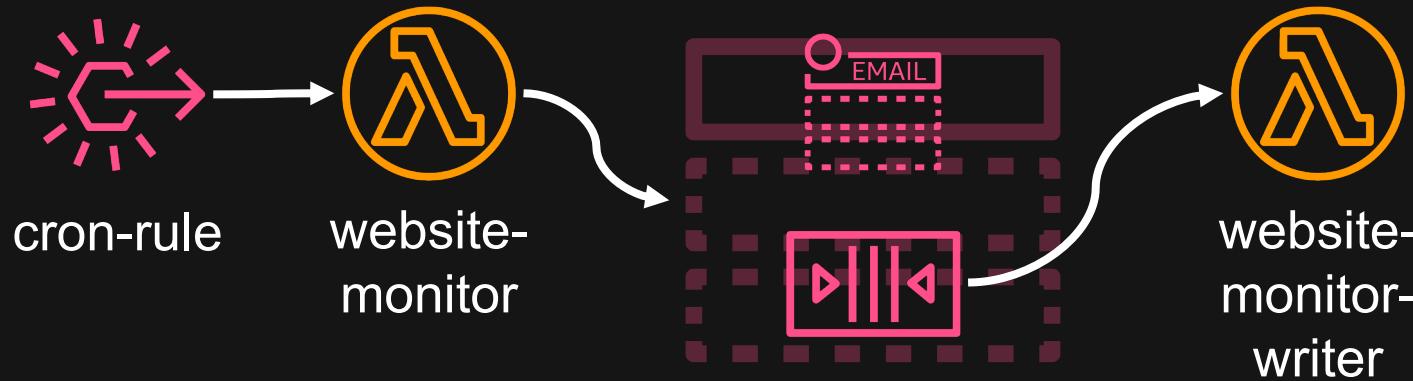


Amazon Simple
Queue Service
(Amazon SQS)

LAB 4 – Complete Architecture



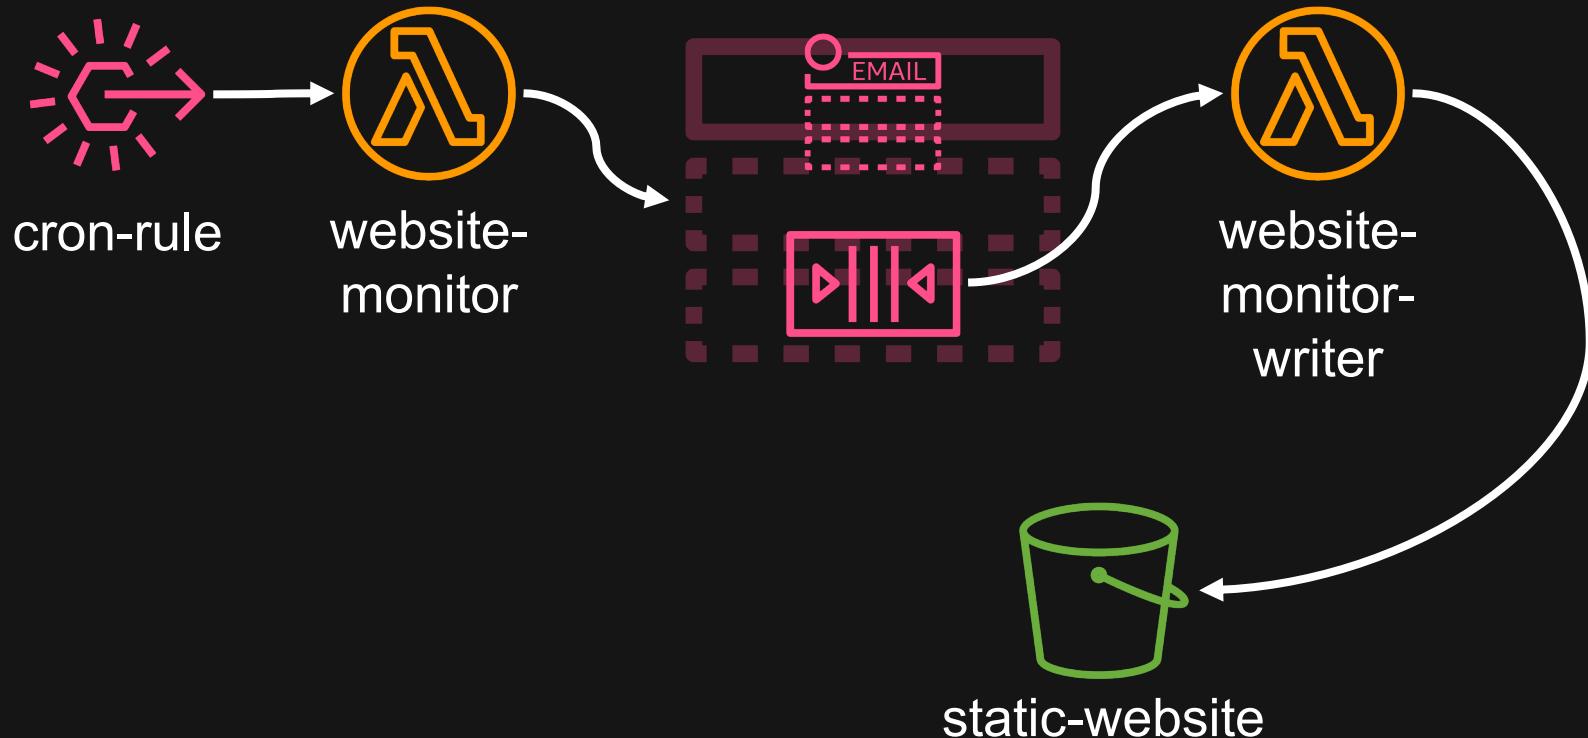
LAB 4 – Complete Architecture



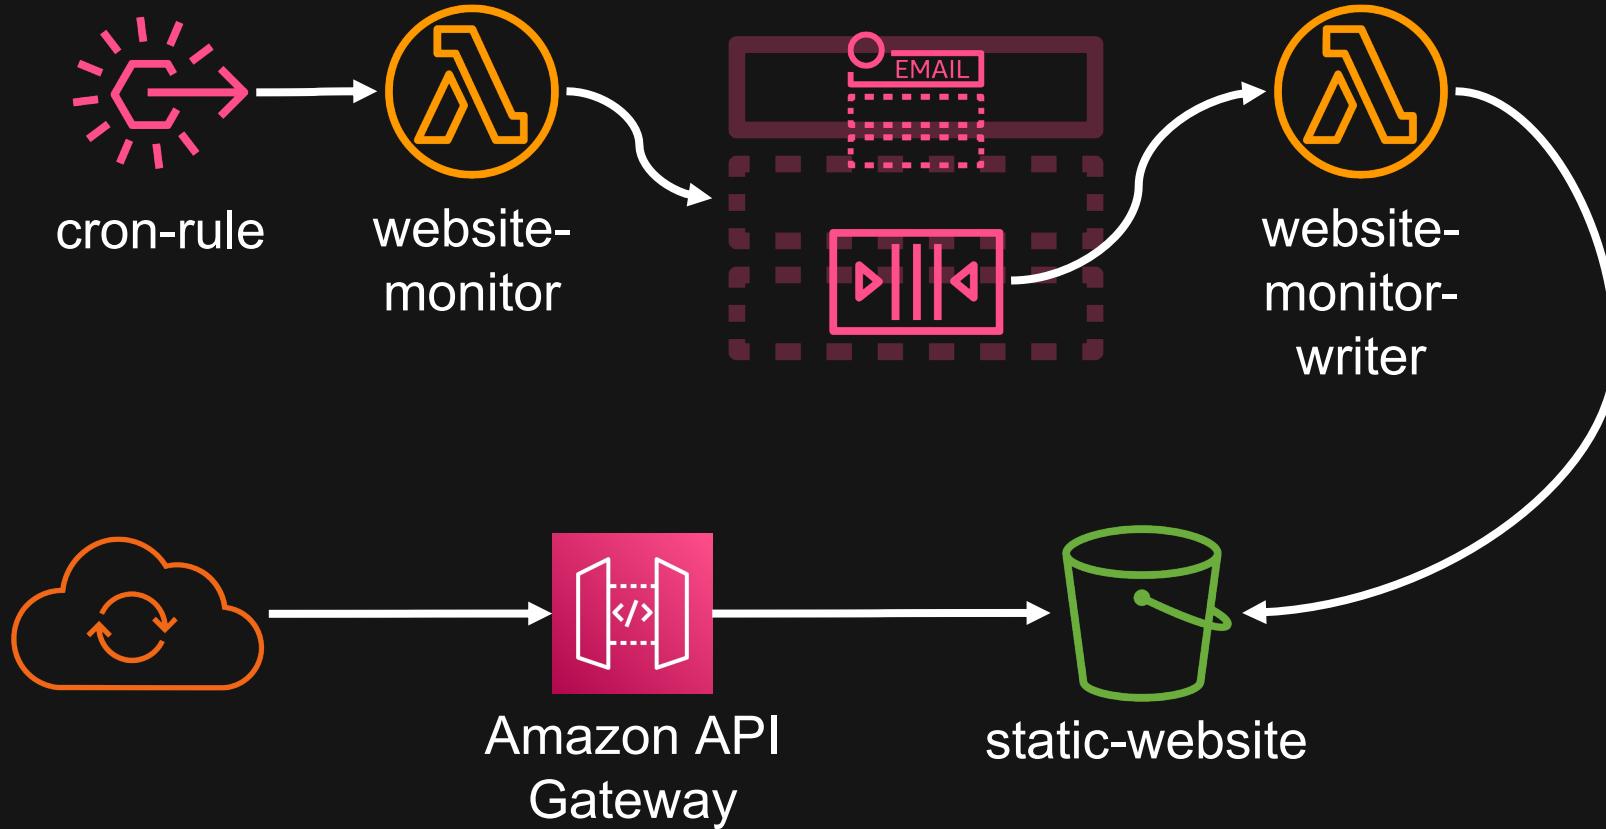
LAB 4 – Complete Architecture



LAB 4 – Complete Architecture



LAB 4 – Complete Architecture





Lab 4: Serverless Capstone Project

INE Cloud Series

An INE On-Demand Course

LAB 4 – Complete Architecture



Amazon
EventBridge

LAB 4 – Complete Architecture



cron-rule

LAB 4 – Complete Architecture



cron-rule

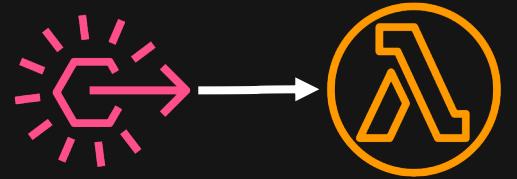


AWS Lambda

LAB 4 – Complete Architecture



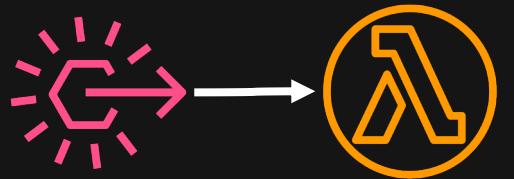
LAB 4 – Complete Architecture



cron-rule

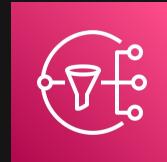
website-
monitor

LAB 4 – Complete Architecture



cron-rule

website-
monitor



Amazon Simple
Notification Service
(Amazon SNS)

LAB 4 – Complete Architecture



LAB 4 – Complete Architecture



LAB 4 – Complete Architecture

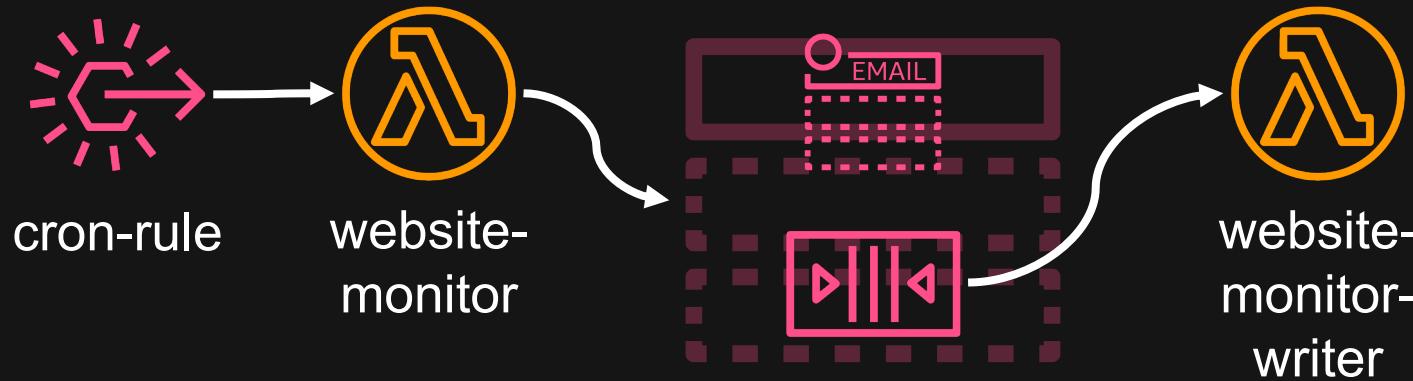


Amazon Simple
Queue Service
(Amazon SQS)

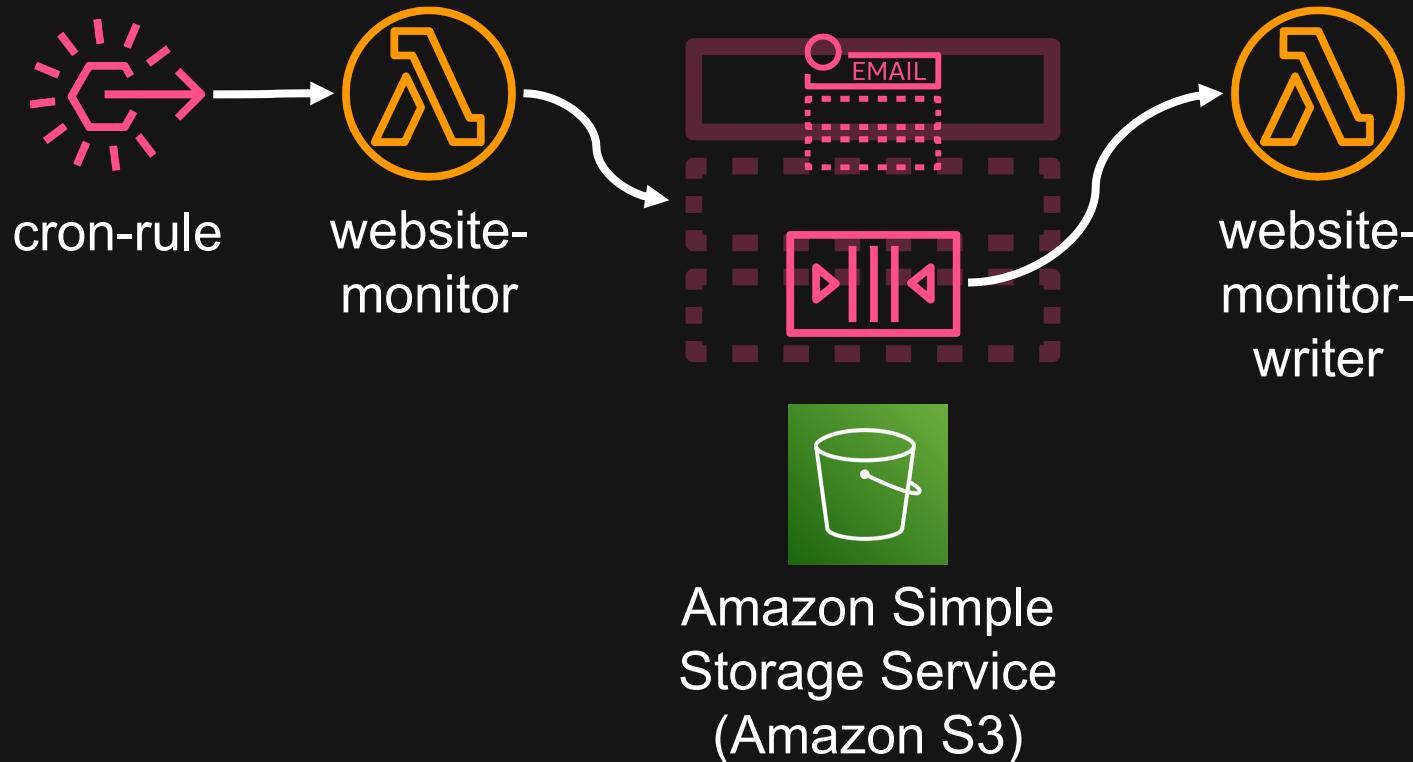
LAB 4 – Complete Architecture



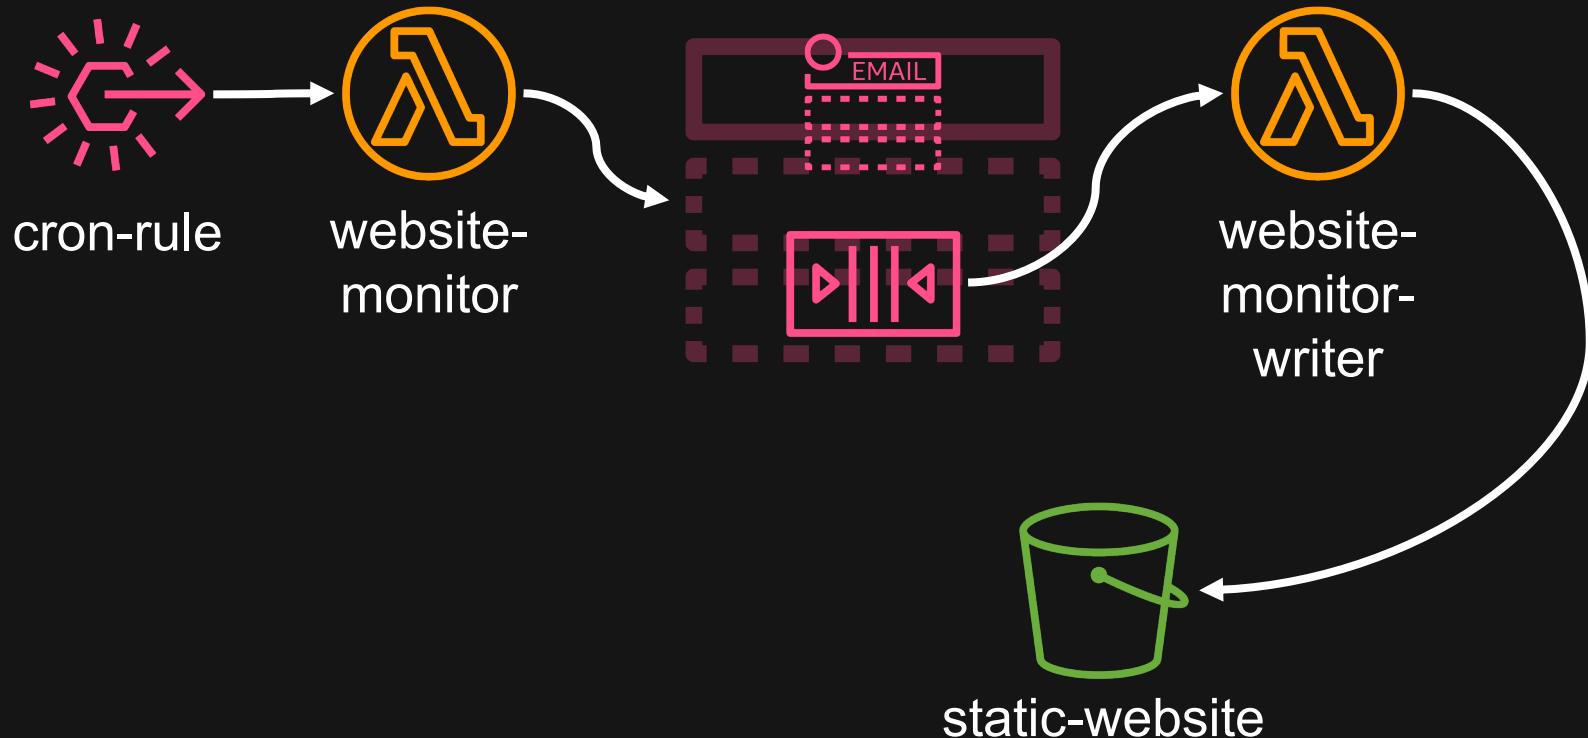
LAB 4 – Complete Architecture



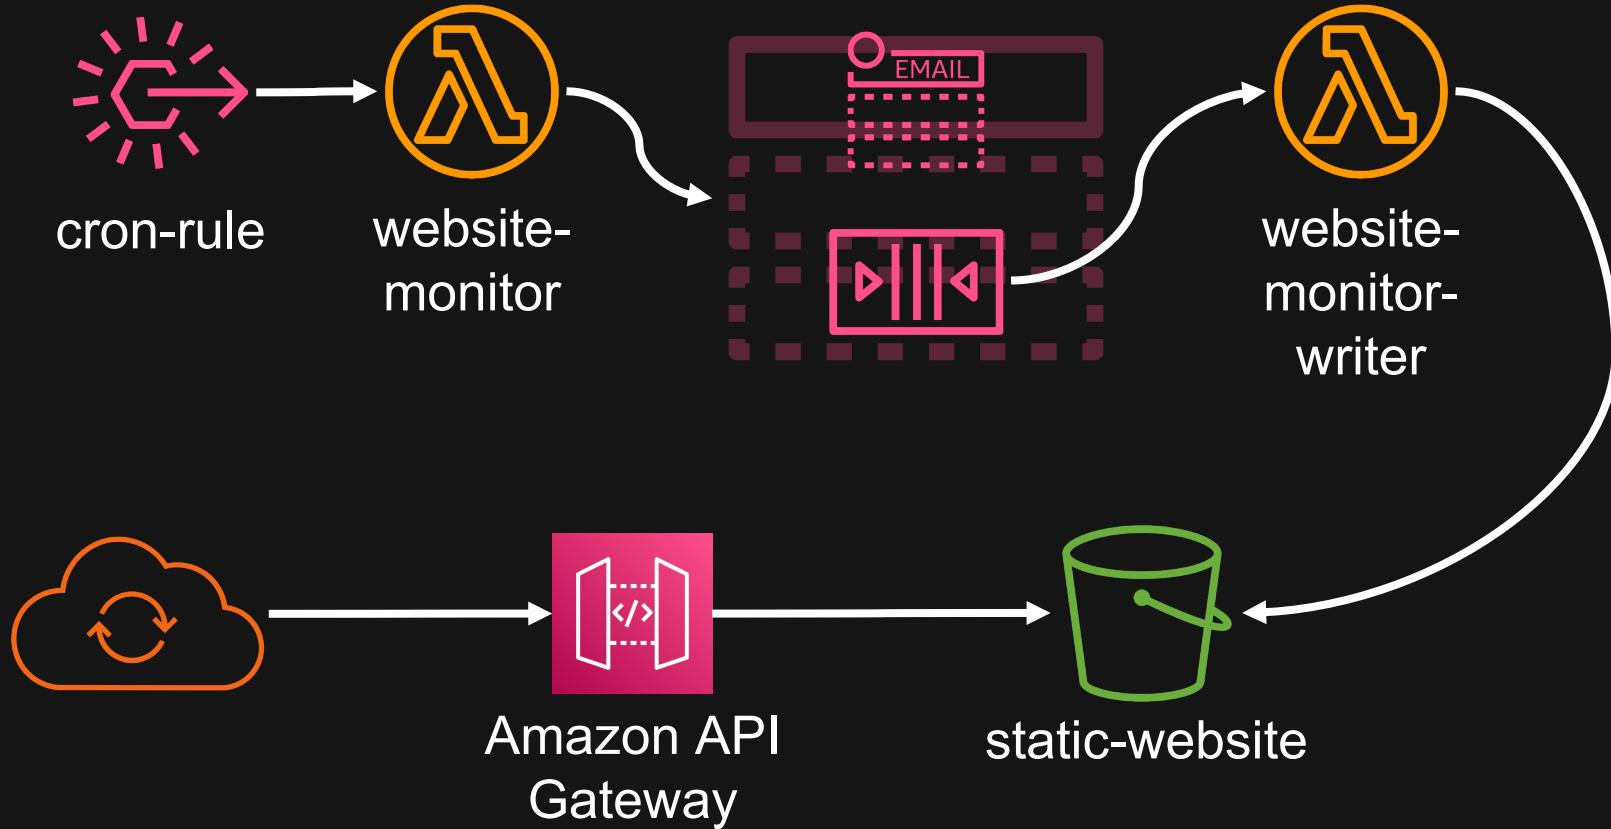
LAB 4 – Complete Architecture



LAB 4 – Complete Architecture



LAB 4 – Complete Architecture





Lab 4: Serverless Capstone Project

Part 8: Code Review

INE Cloud Series

An INE On-Demand Course

